

Causal Clustering



neo4j

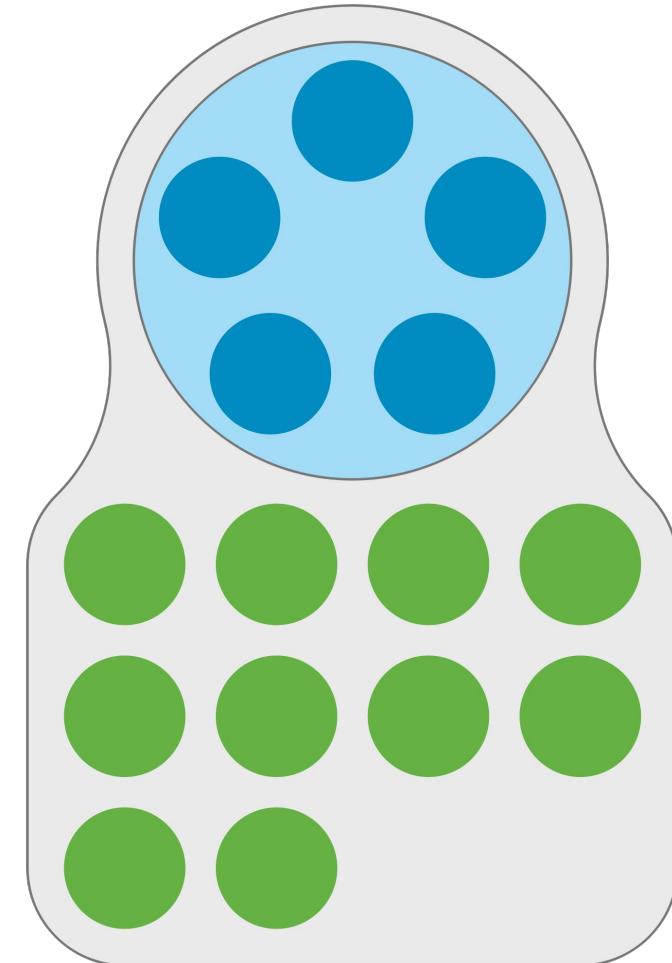


What are we going to do?

- Intro to Causal clustering
- Spin up a Causal cluster
- Cluster Security
- Causal Clustering



Intro to Causal clustering

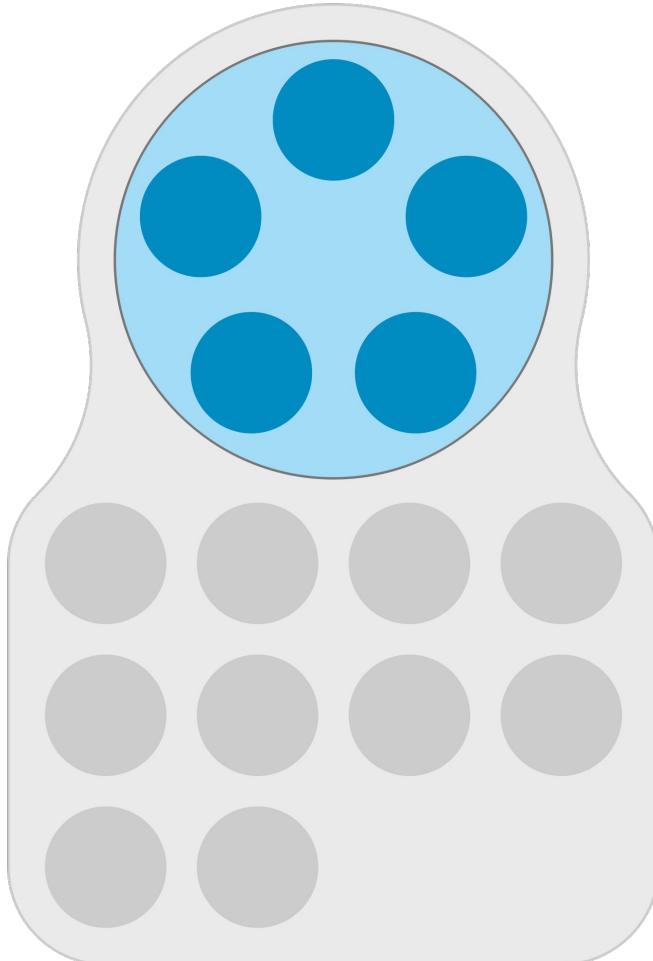


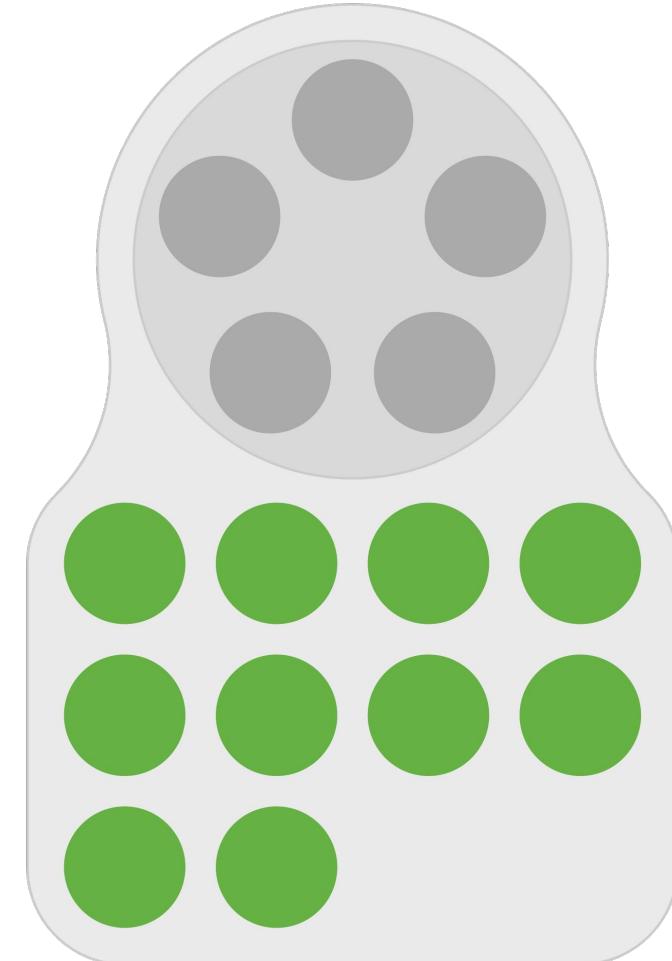
Core

Replica

Core

- Small group of Neo4j databases
- Implements Consensus Commit
- Responsible for data safety





Replica

- For massive query throughput
- Read-only replicas
- Not involved in Consensus Commit
=> less overhead
- Disposable, suitable for auto-scaling



Spinning up a Causal cluster

Download Neo4j Enterprise



Make sure you've got Neo4j Enterprise 3.4.5 on your machine. You can copy this from the USB stick or download from neo4j.com/download

Spin up a cluster



If you can't use the script, you can spin up a cluster manually:

- Make 4 copies of the neo4j directory e.g. core_1 core_2 core_3 edge_1
- In each directory:
 - Replace the config properties from [bit.ly/2qHFnLC](#) in conf/neo4j.conf
 - Start up each of the servers:
`bin/neo4j start`

conf/neo4j.conf



These are the config properties that we've changed:

e.g. *core_1/conf/neo4j.conf*

dbms.mode=CORE

dbms.connector.bolt.listen_address=:7691
dbms.connector.http.listen_address=:7481
dbms.connector.https.listen_address=:7371

causal_clustering.initial_discovery_members=localhost:5001,localhost:5002,localhost:5003
causal_clustering.discovery_listen_address=127.0.0.1:5001
causal_clustering.transaction_listen_address=:6001
causal_clustering.raft_listen_address=:7001



conf/neo4j.conf

These are the config properties that we've changed:

e.g. core_1/conf/neo4j.conf

dbms.mode=CORE

dbms.connector.bolt.listen_address=:7691
dbms.connector.http.listen_address=:7481

dbms.connector.raft.listen_address=:7371

The default mode is 'SINGLE' so if
we want a cluster we need to
explicitly configure the mode.

cluster_members=localhost:5001,localhost:5002,localhost:5003
transaction_listen_address=127.0.0.1:5001
causal_clustering.transaction_listen_address=:6001
causal_clustering.raft_listen_address=:7001

These are the config properties that we've changed:

e.g. *core_1/conf/neo4j.conf*

dbms.mode=CORE

dbms.connector.bolt.listen_address=:7691
dbms.connector.http.listen_address=:7481
dbms.connector.https.listen_address=:7371

causal_clustering.initial_discovery_members=localhost:5001,localhost:5002,localhost:5003
causal_clustering.discovery_listen_address=127.0.0.1:5001
causal_clustering.transaction_listen_address=:6001
causal_clustering.raft_listen_address=:7001

Set the various ports to listen on.
These are used by the Neo4j browser and language drivers.

These are the config properties that we've changed:

e.g. *core_1/conf/neo4j.conf*

```
dbms.mode=CORE
```

```
dbms.connector.bolt.listen_address=:7691  
dbms.connector.http.listen_address=:7481  
dbms.connector.https.listen_address=:7371
```

```
causal_clustering.initial_discovery_members=localhost:5001,localhost:5002,localhost:5003  
causal_clustering.discovery_listen_address=127.0.0.1:5001  
causal_clustering.transaction_listen_address=:6001  
causal_clustering.raft_listen_address=:7001
```

The servers that form the cluster.
This value will be the same on
every cluster member



conf/neo4j.conf

These are the config properties that we've changed:

e.g. *core_1/conf/neo4j.conf*

dbms.mode=CORE

dbms.connector.bolt.listen_address=:7691
dbms.connector.http.listen_address=:7481
dbms.connector.https.listen_address=:7371

causal_clustering.initial_discovery_members=1localhost:5001,localhost:5002,localhost:5003
causal_clustering.discovery_listen_address=127.0.0.1:5001
causal_clustering.transaction_listen_address=:6001
causal_clustering.raft_listen_address=:7001

Ports to listen on for cluster discovery, serving transactions, and consensus commit.

These are the config properties that we've changed:

e.g. *edge_1/conf/neo4j.conf*

dbms.mode=READ_REPLICA

dbms.connector.bolt.listen_address=:7601

dbms.connector.http.listen_address=:7491

dbms.connector.https.listen_address=:7381

causal_clustering.initial_discovery_members=localhost:5001,localhost:5002,localhost:5003

:sysinfo



Navigate to <http://localhost:7481> in your browser and type the following command:

:sysinfo

Core-Edge Cluster Members		
Role	Addresses	Actions
READ_REPLICA	bolt://localhost:7601, http://localhost:7491, https://localhost:7381	Open
LEADER	bolt://localhost:7691, http://localhost:7481, https://localhost:7371	
FOLLOWER	bolt://localhost:7692, http://localhost:7482, https://localhost:7372	Open
FOLLOWER	bolt://localhost:7693, http://localhost:7483, https://localhost:7373	Open

Navigate to LEADER

If the server you just opened doesn't have the LEADER role click on the 'Open' link next to the server which does have that role.

Core-Edge Cluster Members		
Role	Addresses	Actions
READ_REPLICA	bolt://localhost:7601, http://localhost:7491, https://localhost:7381	Open
LEADER	bolt://localhost:7691, http://localhost:7481, https://localhost:7371	
FOLLOWER	bolt://localhost:7692, http://localhost:7482, https://localhost:7372	Open
FOLLOWER	bolt://localhost:7693, http://localhost:7483, https://localhost:7373	Open

Create data



Execute the following query to create some data:

```
UNWIND range(0, 100) AS value  
MERGE (p1:Person {id: value})  
MERGE (p2:Person {id: TOINT(100.0 * rand())})  
MERGE (p1)-[:FRIENDS]->(p2)
```

Create data



Go to one of the servers with the FOLLOWER role by clicking on the 'Open' link next to it. Once you've done that execute the following query:

```
MATCH path = (p:Person)-[:FRIENDS]-(friend)
```

```
RETURN path
```

```
LIMIT 10
```

You can see that the data has been replicated.

Stop the LEADER

Find the server which has the LEADER role:

Core-Edge Cluster Members		
Role	Addresses	Actions
READ_REPLICA	bolt://localhost:7601, http://localhost:7491, https://localhost:7381	Open
LEADER	bolt://localhost:7691, http://localhost:7481, https://localhost:7371	
FOLLOWER	bolt://localhost:7692, http://localhost:7482, https://localhost:7372	Open
FOLLOWER	bolt://localhost:7693, http://localhost:7483, https://localhost:7373	Open

And stop it using the following command:

```
./core_1/bin/neo4j stop
```



Look for the new LEADER

:sysinfo

Core-Edge Cluster Members		
Role	Addresses	Actions
READ_REPLICA	bolt://localhost:7601, http://localhost:7491, https://localhost:7381	Open
FOLLOWER	bolt://localhost:7692, http://localhost:7482, https://localhost:7372	Open
LEADER	bolt://localhost:7693, http://localhost:7483, https://localhost:7373	

Create some more data



Execute the following query to create some more data:

```
UNWIND range(50, 150) AS value  
MERGE (p1:Person {id: value})  
MERGE (p2:Person {id: TOINT(150.0 * rand())})  
MERGE (p1)-[:FRIENDS]->(p2)
```

Bring the old LEADER back again

Now start the old LEADER back up again.

Core-Edge Cluster Members		
Role	Addresses	Actions
READ_REPLICA	bolt://localhost:7601, http://localhost:7491, https://localhost:7381	Open
LEADER	bolt://localhost:7691, http://localhost:7481, https://localhost:7371	
FOLLOWER	bolt://localhost:7692, http://localhost:7482, https://localhost:7372	Open
FOLLOWER	bolt://localhost:7693, http://localhost:7483, https://localhost:7373	Open

```
./core_1/bin/neo4j start
```



Cluster Security



Native Users Cluster

Flexible Authentication Options

Configuring LDAP connector



Use LDAP..

..at this address

..with system account defined

..with this directory structure

```
# Choose LDAP connector as both authentication and authorization provider
dbms.security.auth_provider=ldap

# Configure LDAP connector to point to the AD server
dbms.security.ldap.host=ldap://myactivedirectory.example.com

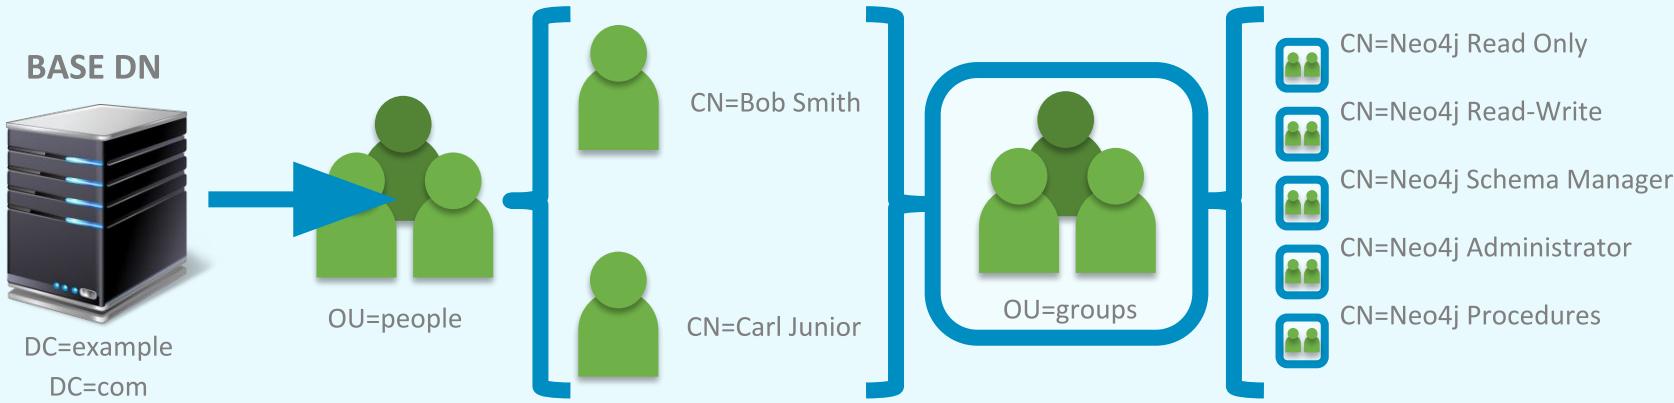
# In case where defined users are not allowed to search for themselves,
# we can specify credentials for user with read access to all users and groups
dbms.security.ldap.authorization.use_system_account=true
dbms.security.ldap.system_username=CN=admin,OU=people,DC=example,DC=com
dbms.security.ldap.system_password=admin-password

# Provide details on user structure within LDAP
dbms.security.ldap.user_dn_template=CN={0},OU=people,DC=example,DC=com
dbms.security.ldap.authorization.user_search_base=OU=people,dc=example,dc=com
dbms.security.ldap.authorization.user_search_filter=(&(objectClass=*)(CN={0}))
dbms.security.ldap.authorization.group_membership_attributes=memberOf
```

. /conf/neo4j.conf

Flexible Authentication Options

LDAP Group to Role Mapping



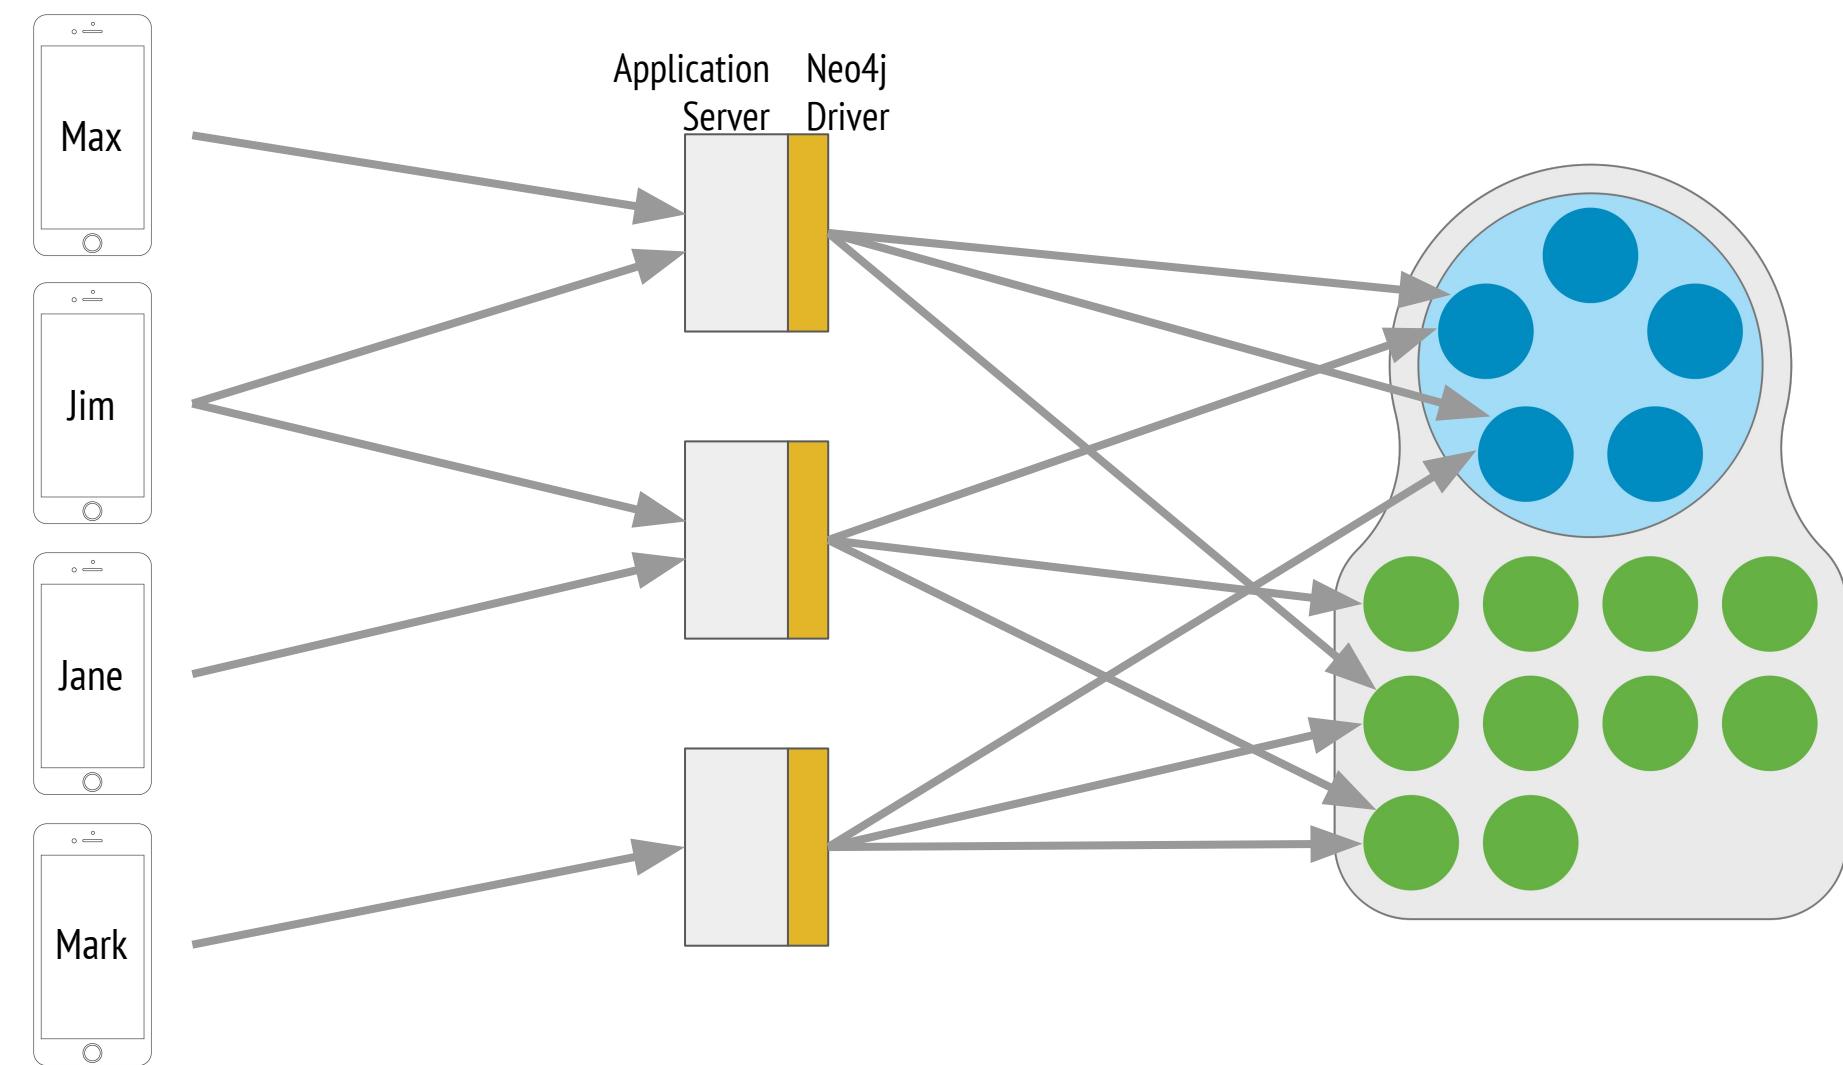
```
# Configure the actual mapping between groups in the LDAP and roles in Neo4j
```

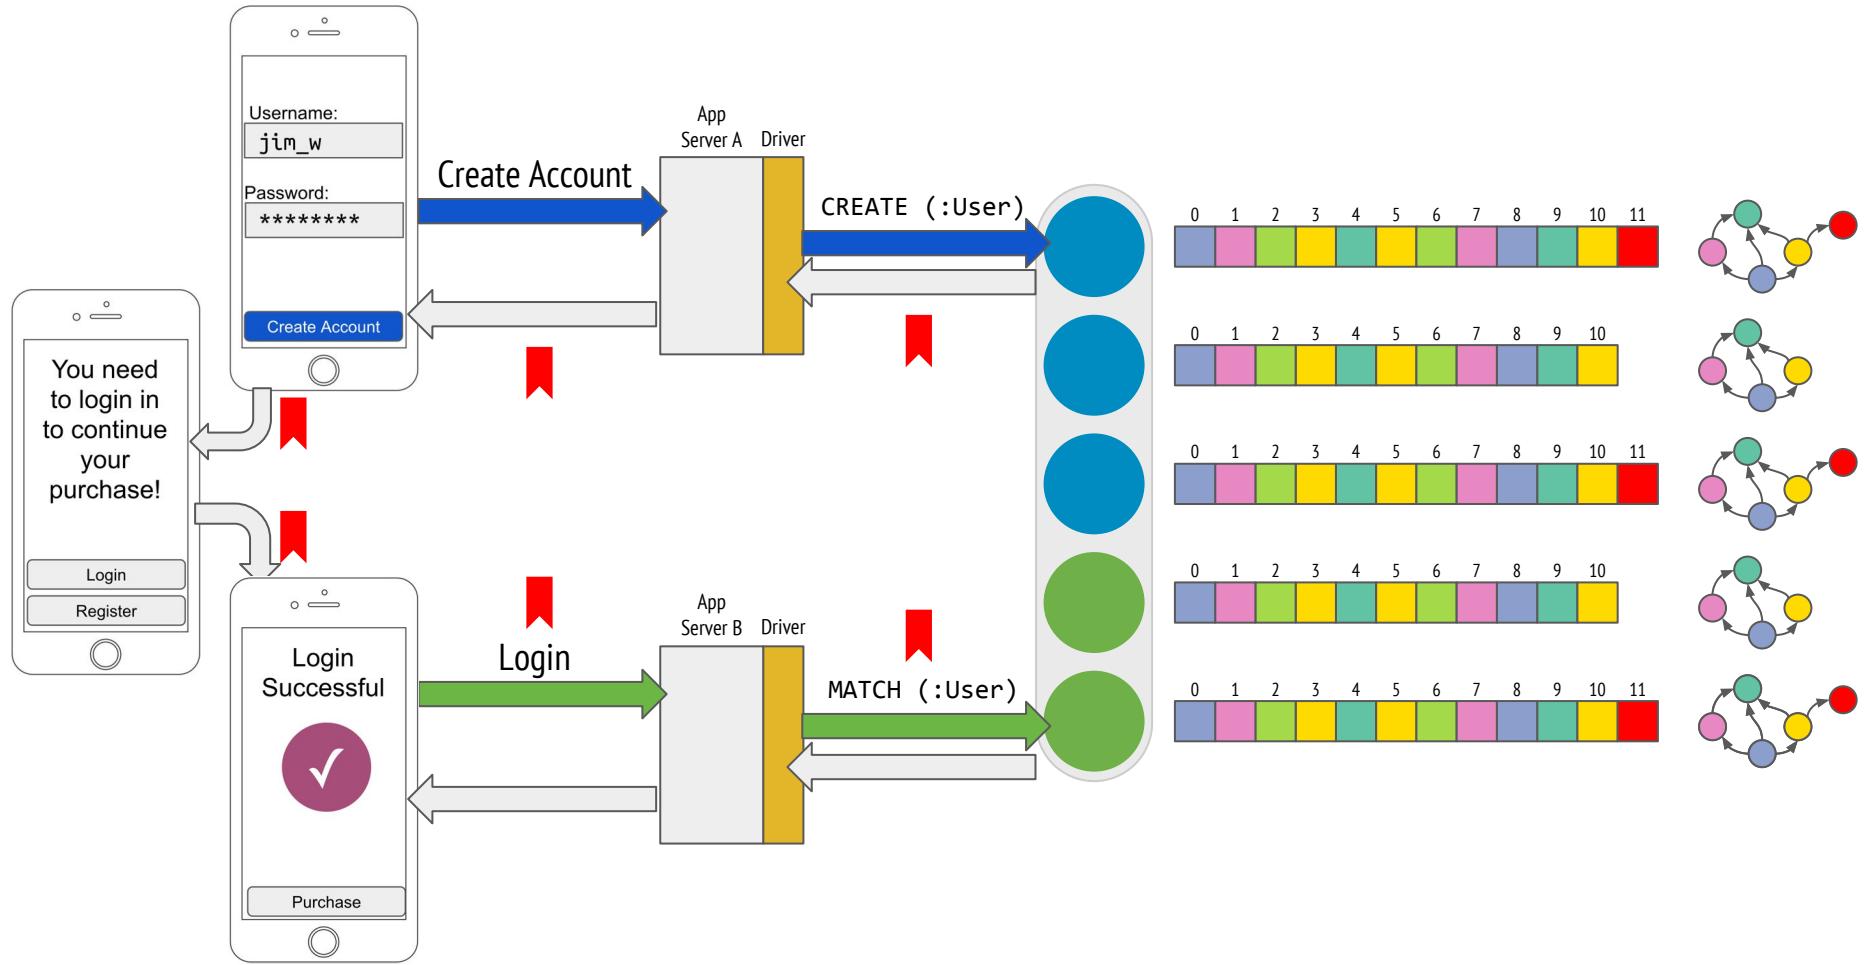
```
dbms.security.ldap.authorization.group_to_role_mapping= \
"CN=Neo4j Read Only,OU=groups,DC=example,DC=com" = reader; \
"CN=Neo4j Read-Write,OU=groups,DC=example,DC=com" = publisher; \
"CN=Neo4j Schema Manager,OU=groups,DC=example,DC=com" = architect; \
"CN=Neo4j Administrator,OU=groups,DC=example,DC=com" = admin; \
"CN=Neo4j Procedures,OU=groups,DC=example,DC=com" = allowed_role
```

[./conf/neo4j.conf](#)



Causal Clustering





Causal Clustering: Read your own writes



```
try ( Driver driver = GraphDatabase.driver( "bolt+routing://localhost:7693",
                                             AuthTokens.basic( "neo4j", "abc" ) ) ) {
    try ( Session session = driver.session() ) {

        session.writeTransaction( (tx) -> {
            return tx.run( "CREATE (person:Person {name: {name}, title: {title}})",
                          parameters( "name", "Webber", "title", "Dr" ) );
        } );

        String title = session.readTransaction( (tx) -> {
            return tx.run( "MATCH (person:Person {name: {name}}) RETURN person.title",
                          parameters( "name", "Webber" ) ).single().get( 0 ).asString();
        } );

        System.out.println( "Webber's title is " + title );
    }
}
```

Causal Clustering: Read your own writes



```
try ( Driver driver = GraphDatabase.driver( "bolt+routing://localhost:7693",
                                             AuthTokens.basic( "neo4j", "abc" ) ) ) {
```

```
    try ( Session session = driver.session() ) {
```

Application holds one global instance of the driver. The Driver is thread safe.

```
        -> {
            Person:Person {name: {name}, title: {title}})",
            "Webber", "title", "Dr" ) );
```

```
        String title = session.readTransaction( (tx) -> {
            return tx.run( "MATCH (person:Person {name: {name}}) RETURN person.title",
                          parameters( "name", "Webber" ) ).single().get( 0 ).asString();
        } );
```

```
        System.out.println( "Webber's title is " + title );
```

```
}
```

Causal Clustering: Read your own writes

```
try ( Driver driver = GraphDatabase.driver( "bolt+routing://localhost:7693",
                                             AuthTokens.basic( "neo4j", "abc" ) ) ) {
    try ( Session session = driver.session() ) {
        String title = session.readTransaction( (tx) -> {
            return tx.run( "MATCH (person:Person {name: $name})
                           SET person.title = $title
                           RETURN person.title"
                         .parameters( "name", "Webber" )
                         .parameters( "title", "Dr" ) );
        } );
        System.out.println( "Webber's title is " + title );
    }
}
```

Application holds one global instance of the driver. The Driver is thread safe.

Bolt routing URI allows access to automatic cluster member routing behaviour.

Causal Clustering: Read your own writes



```
try ( Driver driver = GraphDatabase.driver( "bolt+routing://localhost:7693",
                                             AuthTokens.basic( "neo4j", "abc" ) ) ) {
    try ( Session session = driver.session() ) {
        session.writeTransaction( tx) -> {
            return tx.run("CREATE (person:Person {name: {name}, title: {title}})",
                         "name", "Webber", "title", "Dr" );
        } );
        String title = session.readTransaction( tx) -> {
            return tx.run("MATCH (person:Person {name: {name}}) RETURN person.title",
                         parameters( "name", "Webber" ) ).single().get( 0 ).asString();
        } );
        System.out.println( "Webber's title is " + title );
    }
}
```

The Session is a logical context
for transactional units of work. A
Session is not thread safe.

Causal Clustering: Read your own writes



```
try ( Driver driver = GraphDatabase.driver( "bolt+routing://localhost:7693",
                                             AuthTokens.basic( "neo4j", "abc" ) ) ) {
    try ( Session session = driver.session() ) {
        session.writeTransaction( (tx) -> {
            return tx.run( "CREATE (person:Person {name: {name}, title: {title}})",
                          parameters( "name", "Webber", "title", "Dr" ) );
        } );
        String title = session.readTransaction( (tx) -> {
            return tx.run( "MATCH (person:Person {name: {name}}) RETURN person.title",
                          parameters( "name", "Webber" ) ).single().get( 0 ).asString();
        } );
        System.out.println( "Title is " + title );
    }
}
```

A unit of work is encapsulated within a *Transaction Function*. If execution fails, retries are attempted automatically.

Causal Clustering: Read your own writes



```
try ( Driver driver = GraphDatabase.driver("bolt://localhost:7693",  
    neo4j, "abc" ) ) {  
    try ( Session session = driver.session() ) {  
        session.writeTransaction( () -> {  
            tx.run( "CREATE (p:Person {name: $name}) SET p.title = $title",  
                parameters( "name", "Webber" ),  
                parameters( "title", "Dr" ) );  
        } );  
  
        String title = session.readTransaction( (tx) -> {  
            return tx.run( "MATCH (person:Person {name: $name}) RETURN person.title",  
                parameters( "name", "Webber" ) ).single().get( 0 ).asString();  
        } );  
  
        System.out.println( "Webber's title is " + title );  
    }  
}
```

Cypher results should be fully consumed within a transaction function but values can be returned.

Causal Clustering: Read your own writes



```
try ( Driver driver = GraphDatabase.driver( "bolt+routing://localhost:7693",
                                             AuthTokens.basic( "neo4j", "abc" ) ) ) {
    try ( Session session = driver.session() ) {

        session.writeTransaction( (tx) -> {
            return tx.run( "CREATE (person:Person {name: {name}, title: {title}})",
                          parameters( "name", "Webber", "title", "Dr" ) );
        } );

        String title = session.readTransaction( (tx) -> {
            return tx.run( "MATCH (person:Person {name: {name}}) RETURN person.title",
                          parameters( "name", "Webber" ) ).single().get( 0 ).asString();
        } );

        System.out.println( "Webber's title is " + title );
    }
}
```



Causal Clustering Demo

End of Module Core-Edge Clustering

Questions?



What are we going to do?



Now we're ready to move from our single server to a clustered server setup.

Step 1: Stop the database

Step 2: Create a backup

Step 3: Restore the backup on each of the servers

Step 4: Start the cluster

What are we going to do?



Run a query on the leader and watch it propagate

Demo of config on separate servers?

Gist that we can copy/paste from

Docker instructions? Steal Jonas' script (try it out)

Download cluster bootstrapping script



Your USB stick contains a file called `neo4j-cluster-rsync-auth-master.zip` which we'll use to create a local cluster. Copy this file onto your machine and unpack it.

If you don't have a USB stick you can download it from here:

<https://github.com/craigtaverner/neo4j-cluster-rsync-auth>



Spin up a cluster

Set the following environment variables:

```
export NEO4J_VERSION="3.1.4"  
export NEO4J_PASSWORD="abc"  
export NUMBER_CORES=3  
export NUMBER_EDGES=1
```

Run the installation script:

```
./install.sh
```



Spin up a cluster

Set the following environment variables:

```
export NEO4J_VERSION="3.1.4"  
export NEO4J_PASSWORD="abc"  
export NUMBER_CORES=3  
export NUMBER_EDGES=1
```

Run the installation script:

```
./install.sh
```

If you can't use the script, you can spin up a cluster manually:

- Make 3 copies of the neo4j directory e.g.
`core_1 core_2 core_3 edge_1`
- In each directory:
 - Copy config from
<http://bit.ly/gcsf2016-prod> into
`conf/neo4j.conf`
 - Start up the server:
`bin/neo4j start`



Native Users Cluster Demo

Flexible Authentication Options

Configuring LDAP connector



Use LDAP..

..at this address

..with system account defined

..with this directory structure

```
# Choose LDAP connector as both authentication and authorization provider
dbms.security.auth_provider=ldap

# Configure LDAP connector to point to the AD server
dbms.security.ldap.host=ldap://myactivedirectory.example.com

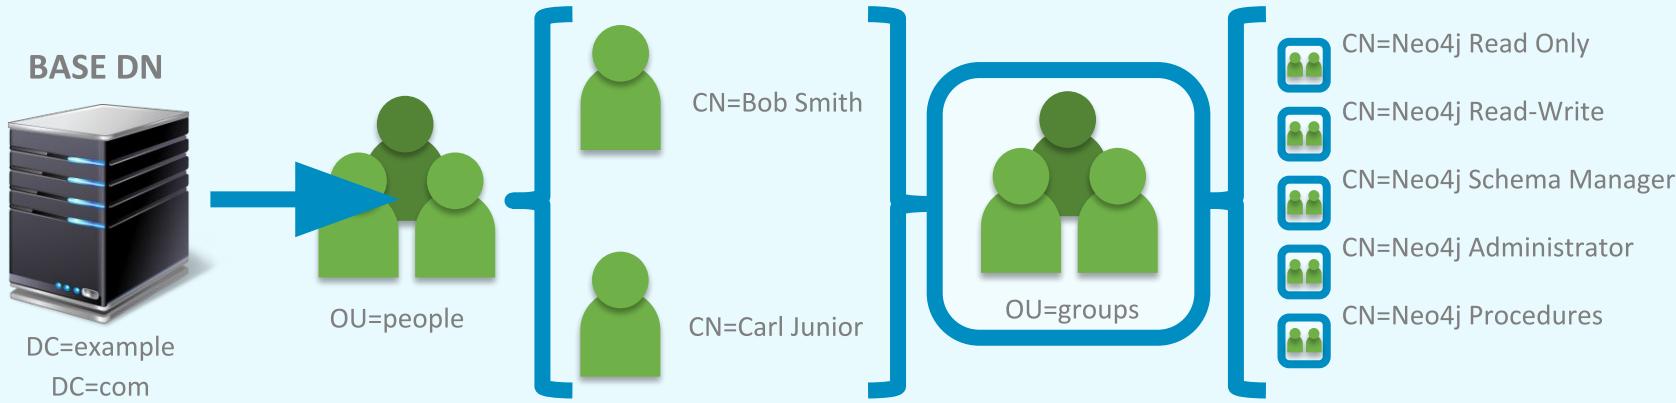
# In case where defined users are not allowed to search for themselves,
# we can specify credentials for user with read access to all users and groups
dbms.security.ldap.authorization.use_system_account=true
dbms.security.ldap.system_username=CN=admin,OU=people,DC=example,DC=com
dbms.security.ldap.system_password=admin-password

# Provide details on user structure within LDAP
dbms.security.ldap.user_dn_template=CN={0},OU=people,DC=example,DC=com
dbms.security.ldap.authorization.user_search_base=OU=people,dc=example,dc=com
dbms.security.ldap.authorization.user_search_filter=(&(objectClass=*)(CN={0}))
dbms.security.ldap.authorization.group_membership_attributes=memberOf
```

. /conf/neo4j.conf

Flexible Authentication Options

LDAP Group to Role Mapping



```
# Configure the actual mapping between groups in the LDAP and roles in  
Neo4j
```

```
dbms.security.ldap.authorization.group_to_role_mapping= \  
  "CN=Neo4j Read Only,OU=groups,DC=example,DC=com" = reader; \  
  "CN=Neo4j Read-Write,OU=groups,DC=example,DC=com" = publisher; \  
  "CN=Neo4j Schema Manager,OU=groups,DC=example,DC=com" = architect; \  
  "CN=Neo4j Administrator,OU=groups,DC=example,DC=com" = admin; \  
  "CN=Neo4j Procedures,OU=groups,DC=example,DC=com" = allowed_role
```

[./conf/neo4j.conf](#)



Active Directory Demo

Reconfigure cluster for Active Directory



Run the following script:

```
./configure_provider.sh AD
```

If you can't use the script, you can configure the cluster manually:

- For each server:
 - Copy the contents of neo4j-AD.conf from <http://bit.ly/neo4j-prod-security> into the end of conf/neo4j.conf
 - Restart the server:
`bin/neo4j restart`

Reconfigure cluster for Active Directory



After the cluster has restarted, reload the browser tabs and login again, but now the password will have changed and the admin user is called 'morpheus'.